

SPITFIRE: Toward a Semantic Web of Things

Dennis Pfisterer, Kay Römer, Daniel Bimschas, Oliver Kleine, Richard Mietz, and Cuong Truong, University of Lübeck

*Henning Hasemann, Alexander Kröller, and Max Pagel, Braunschweig University of Technology
Manfred Hauswirth, Marcel Karnstedt, Myriam Leggieri, and Alexandre Passant, DERI, National University of Ireland, Galway*

Ray Richardson, Daysha Consulting

ABSTRACT

The developed world is awash with sensors. However, they are typically locked into unimodal closed systems. To unleash their full potential, access to sensors should be opened such that their data and services can be integrated with data and services available in other information systems, facilitating novel applications and services that are based on the state of the real world. We describe our vision and architecture of a Semantic Web of Things: a service infrastructure that makes the deployment and use of semantic applications involving Internet-connected sensors almost as easy as building, searching, and reading a web page today.

INTRODUCTION

Sensors are ubiquitous in infrastructures, appliances, mobile phones, and wireless sensor networks. Their widespread deployment represents significant financial investment and technical achievement, and the data they deliver is capable of supporting an almost unlimited set of high value proposition applications. This is a powerful and profitable confluence of need, capability, and economic opportunity — yet the true potential of sensor technology is massively under-exploited.

A central problem hampering success is that sensors are typically locked into unimodal closed systems. For example, motion detection sensors in a building may be exclusively controlled by the intrusion detection system. Yet the information they provide could be used by many other applications, for example, placing empty buildings into an energy-conserving sleep mode or locating empty meeting rooms. Unlocking valuable sensor data from closed systems has the potential to revolutionize how we live. To realize this potential, a service infrastructure is needed to connect sensors to the Internet and publish their outputs in well-understood machine-pro-

cessible formats on the web, thus making them accessible and usable on a large scale under controlled access.

So far, the sensor world and the web world have been largely disconnected, requiring the human in the loop to find, integrate, and use information and services from both worlds in a meaningful way. Publishing sensor-related data on the web would help people find relevant information by directly accessing sensor data, that is, by directly observing the real world, integrated with related information from the web. Already today, smart phone applications such as CenceMe exist that infer the activity of the person wearing the phone from sensor data and publish this on the web. Another example is energy consumption sensors such as Ploggs that end users can install in their homes to measure energy consumption of appliances, for example, to compare their energy consumption with that of other, similar households to identify opportunities for saving energy. To do this easily, with open interfaces and data formats, and on a large scale, technologies from the web need to be customized for and integrated with their relevant counterparts on the Internet of Things (IoT). This means that application experts who are able to publish web pages today should have the same easy-to-use technologies at hand to publish sensor descriptions and sensor data, and make use of sensor outputs without requiring deep knowledge of embedded computing. In particular, we believe that users are primarily interested in real-world entities (things, places, and people) and their high-level states (*empty, free, sitting, walking*, etc.) rather than in individual sensors and their raw output data. Therefore, the infrastructure has to provide appropriate abstractions to map sensors and their raw output to real-world entities with high-level states.

Real-world entities are rarely useful when considered in isolation; the ability to put multiple entities into a common semantic context is needed. For example, we want to reason about

This work has been partially supported by the European Union under contract number ICT-2009-258885 (SPIT-FIRE), by Science Foundation Ireland under Grant No. SFI/09/CE/I1380 (Lion-2), and by the Federal Ministry of Education and Research of Germany (Förder Kennzeichen 01BK0905, Real-World G-Lab).

rooms being in the same building, belonging to the same company, with nearby parking spots. This requires a machine-readable representation of world knowledge and appropriate reasoning capabilities. Furthermore, this representation needs to be unified — while most sensor data published so far on the web relies on heterogeneous data models and serializations. In addition to discovery and query facilities on static properties of those machine-readable representations of sensors and real-world entities, specialized search approaches to support queries on the dynamically changing state of sensors or entities consisting of many sensors (possibly integrated with static data) will be required (e.g., which rooms in a building are currently occupied).

There are efforts to realize a Semantic Sensor Web, including the SENSEI, SemSorGrid4Env, Exalted, and 52 North projects, as well as work by the Kno.e.sis Center, CSIRO, and the Spanish Meteorological Agency. Most notably, the Open Geospatial Consortium's (OGC) Sensor Web Enablement (SWE) [1] project builds a framework to publish and access sensor data using XML-based protocols and application programming interfaces (APIs).

The choice of XML, however, ties SWE to system-specific schemas, providing neither semantic interoperability nor a basis for reasoning. This problem is the focus of the Semantic Sensor Web [2], which proposes annotating sensor data with semantic meta-data, the meaning of which is machine-understandable through vocabulary definitions (i.e., an ontology). By annotating sensor-related features such as the network, deployment, data formats, and so on, it becomes possible to automate further tasks (e.g., deployment, maintenance, and integration).

However, these efforts have limitations, which we are addressing in SPITFIRE: There is no general-purpose approach compatible with the growing body of semantic world knowledge available as linked open data (LOD) on the web; existing efforts are either too sensor-centric or too knowledge-centric; that is, they do not provide comprehensive, integrated abstractions for things, their high-level states, and how they are linked to sensors; and a number of important services are missing in existing efforts, notably support for semi-automatically creating linked data representations of sensors and things, as well as efficient search for things based on their current states.

SPITFIRE addresses these limitations by providing:

- Vocabularies to integrate descriptions of sensors and things with the LOD cloud
- Semantic entities as an abstraction for things with high-level states inferred from embedded sensors
- Semi-automatic generation of semantic sensor descriptions
- Efficient search for sensors and things based on their current states

In addition, SPITFIRE integrates these ingredients into a unified service infrastructure to ease adoption of the Semantic Web of Things for end users and developers. On top of this infrastructure, applications are assembled by issuing search requests for matching (real or

aggregated) sensor services and invoking found services directly.

This article proceeds with the description of an exemplary use case to illustrate the state of the art with respect to integration of sensors into the web upon which SPITFIRE builds, followed by a description of the novel contributions of SPITFIRE and a brief discussion of an existing operational prototype.

USE CASE AND REQUIREMENTS

Due to an emergency, a traveling salesman drives to his company headquarters to hold an ad hoc meeting. For that, he must find a free room in the headquarters, dispersed over a large area. After that, he informs his colleagues and searches for a parking spot close to the building.

Imagine that sensors, which are connected to the Internet, measure the state of real-world entities such as meeting rooms and parking spots. Internet connectivity not only requires network-level integration (IP), but also application-level integration to enable structured access to sensor data. To enable automatic reasoning about sensors (e.g., finding free parking spots close to a meeting room), these sensors, their output, and their embedding into the real world must be *described in a machine-readable format* that is compatible with data formats used to describe existing world knowledge in the web. Not only must the syntax and semantics of such a description be defined, but efficient mechanisms to *annotate newly deployed sensors* with appropriate descriptions are required.

Users are primarily interested in *real-world entities* (e.g., “meeting room”) and their *high-level states* (e.g., “room occupied”) rather than sensors (e.g., “sensor 536”) and their raw output (e.g., “motion detected at time T ”). Therefore, appropriate mechanisms to establish an explicit mapping of sets of sensors to the real-world entities they are monitoring (e.g., all motion detection sensors in a certain room) must be provided. Furthermore, the raw output of these sensors (e.g., motion detection events) has to be mapped to a high-level state (e.g., “room occupied”). Often, this involves fusing the outputs of multiple sensors (e.g., multiple motion sensors are needed to cover a large room) or even scheduling sensors for energy efficiency (e.g., only one of two available battery-powered motion sensors is required to cover a smaller room).

Finally, the user wants to *search for real-world entities by their current state* (e.g., empty meeting rooms). Often, such search requests refer not only to the output of sensors, but also to further machine-readable information that is available elsewhere in the web (e.g., company maps, meeting schedules, calendars). The search engine needs to integrate these different static and dynamic data sources in a seamless way.

SEMANTIC SENSOR WEB: STATE OF THE ART

Realizing the above use case on an Internet scale requires:

- That the sensors are connected to the Internet

To enable automatic reasoning about sensors, these sensors, their output, and their embedding into the real world must be described in a machine-readable format that is compatible with data formats used to describe existing world knowledge in the web.

Integrating resource-constrained sensors into the Internet is difficult since ubiquitously deployed Internet protocols such as HTTP, TCP, or even IP are too complex and resource-demanding. To achieve integration, lightweight alternatives are required that can easily be converted from/to Internet protocols.

- That machines can discover and understand the semantics of the data returned by the sensors
- A technique to find the sensors that could provide the relevant data

This section briefly discusses the state of the art in relation to this with a focus on Internet-scale web-based technologies upon which SPITFIRE builds, employing the use case as an example.

CONNECTING SENSORS TO THE INTERNET AND THE WEB

Integrating resource-constrained sensors into the Internet is difficult since ubiquitously deployed Internet protocols such as HTTP, TCP, or even IP are too complex and resource-demanding. To achieve integration, lightweight alternatives are required that can easily be converted from/to Internet protocols.

Only recently, two such alternatives are gaining momentum: 6LoWPAN and the Constrained Application Protocol (CoAP). 6LoWPAN [3] is a lightweight IPv6 adaptation layer allowing sensors to exchange IPv6 packets with the Internet. Currently, only UDP is specified as TCP is considered too resource-consuming. CoAP [4] is a draft by the Internet Engineering Task Force's (IETF's) CoRE working group, which deals with Constrained RESTful Environments. It provides a lightweight alternative to HTTP using a binary representation and a subset of HTTP's methods (GET, PUT, POST, and DELETE). In addition, CoAP provides some transport reliability using acknowledgments and retransmissions. For seamless integration, reverse proxies may convert 6LoWPAN to IPv6 and CoAP/UDP to HTTP/TCP so that sensor data can be accessed using these omnipresent protocols. Also, Internet-based clients could directly use CoAP on top of UDP.

6LoWPAN in combination with CoAP allows sensors to be queried from the Internet as they can provide so-called RESTful web services. Those are services following the Web's REST (REpresentational State Transfer) principles [5]. Resources (e.g., sensors) are addressed using standard URIs and data can be returned in different representations (e.g., HTML or Resource Description Framework [RDF]) using HTTP/CoAP content negotiation.

RESTful services are queried and manipulated using the aforementioned four HTTP/CoAP methods. For instance, an application could query the state of a sensor by sending a GET request to the sensor (e.g., `coap://ipv6-address-or-dns-name/room-sensor`). The sensor replies with its value encoded in a possibly proprietary encoding (e.g., in plain text: *occupied* or any format the sensor supports). For an exhaustive discussion of 6LoWPAN, CoAP, and RESTful services, we refer the reader to [6].

To discover the services hosted on a CoAP server, the CoRE link format specification defines how web linking described in IETF RFC 5988 is used by CoAP servers. Clients use a well-known URI (`/well-known/core`) to retrieve a list of resources. For instance, the room sensor device could return `</room-sensor>;ct=0;rt="ex:RoomSensor"` to indicate that the resource

/room-sensor returns the content type *text/plain* (indicated by `ct=0`) and that the resource type is *ex:RoomSensor*. The latter is a concept from an ontology (e.g., the World Wide Web Consortium [W3C] SSN-XG sensor ontology to which some of the authors contributed) as described in the following section. Note that concepts appearing in different ontologies can be automatically mapped.¹

RESTful services (i.e., the operations provided, their parameters and return values) can be described using, for example, Web Service Definition Language (WSDL) version 2.0. For example, RESTful versions of OGC's Sensor Observation Services have been proposed and are currently under consideration by the Sensor Web Enablement group.

LINKED SENSOR DATA

The integration of sensors into the Internet using CoAP/HTTP already enables many applications in which developers query and process data provided by a known set of sensors. However, such manual integration does not scale. What is required is a "machine-understandable" description of sensors and the data they produce. Semantic Web [7] technologies fulfill this requirement as they enable machines to understand, process, and interlink data using structured descriptions of resource and Linked Open Data as the framework makes this integration both immediate and meaningful through the inclusion of semantic links into a resource's machine-readable description.

The predominant technique for *machine-readable representations of knowledge* on the Web is the RDF, which represents knowledge as (*subject, predicate, object*)-triples (e.g., *Sensor3 is-in ParkingSpot41* or *ParkingSpot41 is-in Berlin*). A set of triples forms a graph where subjects and objects are vertices and predicates are edges. From the graph formed by these two triples, one can infer that *Sensor3* is in Berlin by exploiting the knowledge (contained in so-called ontologies) that *is-in* is a transitive property. Such knowledge is often expressed using Web Ontology Language (OWL), one of the main languages (with RDF schema) to define ontologies on the web.

It is imperative to use non-ambiguous identifiers for subjects, predicates, and objects to guarantee uniqueness on an Internet scale, which is achieved by encoding them as URIs. The above triple could be expressed as follows:

- A subject (`<http://example.com/sensors/sensor3>`)
- A predicate (`<http://www.loa-cnr.it/ontologies/DUL.owl#hasLocation>`)
- An object (`<http://example.com/parkingSpot/spot41>`)

The linked data model does not enforce special URIs, but encourages the use of widely used URIs so that a densely interlinked graph emerges. *Ontologies* play an important role in defining the URIs for a specific application domain and their relation to each other as they "standardize" agreed conceptual knowledge. For example, an ontology could define a generic sensor (e.g., `http://purl.oclc.org/NET/ssnx/ssn#Sensor`), an occupancy detection sensor

¹ <http://www.w3.org/DesignIssues/RDF-XML>

(e.g., `http://example.com/ontology/spitfire.owl#Occupancy`), and define that occupancy sensor is a subclass of sensor, which creates a relation between the two URIs. Semantic search engines queried for sensors at a certain location could therefore specifically return information on *occupancy detection sensors*. The CoAP link-format we use (RFC 5988) allows URIs to be specified, eventually pointing to semantic definitions (i.e., support for semantic annotation of links inside a sensor network). Additionally, to make these semantic descriptions available on the web, we could imagine annotating pages describing sensors using RDF or SA-REST so that the same document can be used for humans and machines.

SEARCH FOR SENSORS

Assuming that sensors are described by such RDF triples, a search service can find sensors based on meta-data such as sensor type, location, or accuracy. For instance, applications could ask for parking spots in Berlin to calculate the city's availability of car parking places. Such queries can be expressed in SPARQL, and the aforementioned question could be answered using the (simplified) SPARQL query in Fig. 1. In the query, question marks indicate variables (e.g., “node” and “spot”), while “spots” is an aggregate value.

The variables in a SPARQL query are matched against triples in databases (triple stores) and bound to the matching fields in the matching triples. That is, the query finds subjects that are sensors observing occupancy that are located in a spot that is a parking spot located in Berlin. There are a number of existing efforts to support semantic sensor discovery but they are not as comprehensive as SPITFIRE. For example, [8] does not expose linked data, and [9] does not exploit the hierarchical and structured relations which are relevant even for such simple queries as above. To further exploit these annotations, we could also use faceted browsers such as MIT Simile's Exhibit, where facets for identifying parking places could be location, availability, but also static information such as price range.

FROM SEMANTIC SENSOR WEB TO SEMANTIC WEB OF THINGS

The techniques for integration of sensors into the web outlined in the previous section are necessary, but not sufficient to realize a Semantic Web of Things. In particular, semantic descrip-

```

1 | SELECT COUNT(DISTINCT ?node) as ?spots
2 | WHERE {
3 |   ?node a ssn:Sensor ;
4 |     ssn:observes ex:Occupancy ;
5 |     dul:hasLocation ?spot .
6 |   ?spot a ex:ParkingSpot ;
7 |     dul:hasLocation dbpedia:Berlin .
8 | }

```

Figure 1. SPARQL query requesting all occupancy sensors located at parking spots in Berlin.

tions must be integrated with the LOD cloud to support semantic reasoning; semantic descriptions need to be semi-automatically created for sensors and things to allow end users to use this technology at scale; abstractions for things, their high-level states, and integration with sensors are required; and search for things with a given current state is needed. In this section, we show how SPITFIRE addresses these requirements.

ADVANCED SEMANTICS OF SENSORS

Our main contribution in semantics for sensors over projects such as SENSEI and SemSor-Grid4Env are:

- Scalability through avoiding any registry of semantic entities
- Linkage with the LOD cloud not restricted to a predefined set of datasets (as in Kno.e.sis' Real Time Feature Streams and DERI's SensorMasher)
- Generation of semantic annotations from raw data (as in SENSEI) but improved by searching for already existing concepts to reuse and datasets to which to link

To facilitate this, we developed an ontology starting from the alignment (i.e., mapping equivalent concepts in different ontologies, e.g., *room* and *chamber*) of the already existing ontologies such as *Dolce Ultralite*, the W3C Semantic Sensor Network (SSN-XG) ontology (based on OGC's Observations & Measurements and Sensor Model Language), and *Event Model F*, to support cross-domain descriptions of sensor-related data and its context (higher-level events).

As an example, Fig. 2 presents part of an RDF description of an occupancy detection sensor located in a parking spot. The triples state that a particular object is a sensor, measuring occupancy, located in a particular parking spot. The parking spot belongs to a specific company and is located in a given area with a given geographical location (triples following a semicolon have only a predicate and object — the subject is

Over time, more and more sensors are annotated, increasing the chances of our algorithm correctly annotating new sensors. Specifically, if multiple sensors with the same annotation but slightly different output time series exist, a voting mechanism may improve the confidence.

```

1 | <http://sensorportal.tu-braunschweig.de/ontology/spitfire.owl#sensor1>
2 |   rdf:type ssn:Sensor ;
3 |   ssn:observes ex:Occupancy ;
4 |   dul:hasLocation <http://sensorportal.tu-braunschweig.de/parkingSpot/ACMEParkingPlace1> .
5 |
6 | <http://sensorportal.tu-braunschweig.de/parkingSpot/ACMEParkingPlace1>
7 |   rdf:type ex:parkingSpot ;
8 |   dul:isPartOf <http://sensorportal.tu-braunschweig.de/building/plan/companyPlan1> ;
9 |   geo:lat "51" ;
10 |  geo:long "0.4" ;
11 |  dul:hasLocation <http://sensorportal.tu-braunschweig.de/parkingArea/area10> ;
12 |  ssn:attachedSystem <http://sensorportal.tu-braunschweig.de/ontology/spitfire.owl#sensor1> .

```

Figure 2. Exemplary description of a parking spot occupancy detection sensor (RDF serialization in N3).

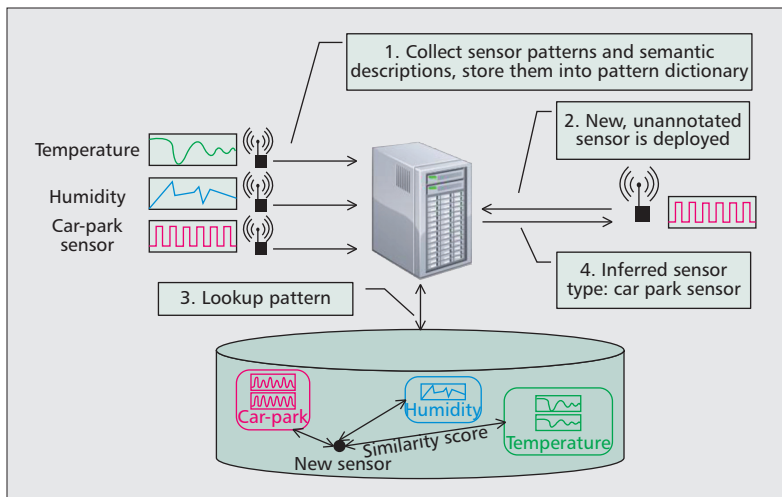


Figure 3. Clusters of annotated sensors and automatic inference of the annotation of a new sensor.

the same as in the previous triple). This shows how a sensor could provide an unambiguous machine-understandable self-description.

SPARQL can be used for search on top of such descriptions, but cannot be used for sensor output directly as this would require new triples to be stored and indexed by a triple store whenever the output of a sensor changes. However, SPARQL-compatible support for this is under active development in SPITFIRE (Continuous Query Evaluation over Linked Streams — CQELS).

A very important feature of our approach is that arbitrary datasets on the LOD can be dynamically linked in order to describe complex real-world processes, and to detect facts and correlations about them. For instance, the above parking spot sensors indirectly also provide information about environmental pollution caused by cars, which can be integrated with additional environmental features from <http://linkedgeodata.org/>, and abnormally high death rates of local population from <http://www4.wiwiw.fu-berlin.de/eurostat/>. All publications about such features can be selected from <http://eprints.rkbexplorer.com/>. Specific information for certain countries such as the United Kingdom is also available as LOD from <http://education.data.gov.uk>.

SEMI-AUTOMATIC CREATION OF SEMANTIC SENSOR DESCRIPTIONS

Besides a language for semantically describing sensors and their output, as well as their embedding into the real world, a mechanism to annotate newly deployed sensors with descriptions in this language is required. The technically simplest way is to let humans provide these, but this approach does not scale as end users typically cannot be expected to provide such semantic descriptions of sensors and their deployment contexts. Therefore, we developed a semi-automatic approach for annotating newly deployed sensors under the hypothesis that sensors with similar semantic descriptions would produce similar output. For example, two motion sensors

deployed in the same room should produce similar time series. Assume that one sensor is manually annotated with its room number, and later, a second sensor is deployed in the same room. If the output of this sensor correlates over some time with the first sensor, we can conclude that they are in the same location and copy the annotation. Based on the strength of the correlation we can compute the confidence of the correctness of the inferred annotation. If there is no sensor with a strong correlation, the user must provide the annotation manually. If there are multiple sensors with different annotations but similar correlation strength, the user must choose among those. That is, initially users must provide the annotations manually to bootstrap the system. Over time, more and more sensors are annotated, increasing the chances of our algorithm to annotate new sensors correctly. Specifically, if multiple sensors with the same annotation but slightly different output time series exist, a voting mechanism may improve the confidence.

The underlying scientific problem is the classification of sensor data streams, which can be solved by means of data clustering as illustrated in Fig. 3. Sensors with the same known annotation form a cluster. When a new sensor with unknown annotation appears, we compute a distance measure to all known clusters, assume the new sensor belongs to the closest cluster, and copy the cluster annotation to the new sensor. However, if the distance to the closest cluster is too large, or there is a tie between multiple clusters, the user is involved as described above. The distance between two sensors is computed by filtering and normalizing the raw sensor output over a given time window, and by computing a distance metric between pairs of sensors. Currently, we use a fuzzy approach for the distance computation, where an output time series of a reference sensor is mapped to fuzzy sets, and the output time series of a newly deployed sensor is matched against those fuzzy sets to compute a similarity score. This algorithm takes into account not only the possible sensor output values, but also their distribution over time. This approach can successfully classify sensors not only according to their type (i.e., measured physical quantity), but according to different measurement units for the same physical quantity (e.g., degrees Celsius vs. degrees Fahrenheit) and according to location (e.g., room where sensors are located).

SEMANTIC ENTITIES: AN ABSTRACTION FOR THINGS

As mentioned earlier, users are interested in real-world entities and their high-level states rather than raw sensor output. We propose semantic entities as a concept to map sensors and their raw output to real-world entities with high-level states. That is, semantic entities provide two orthogonal functions: mapping sensors to real-world entities and computing high-level states from raw output of one or more of those sensors.

This mapping is accomplished using SPARQL queries. For example, to construct a seman-

```

1 SELECT DISTINCT ?node
2 WHERE {
(a) 3 ?node a ssn:Sensor ;
4     dul:hasLocation ex:ACMEParkingPlace1 .
5 }

1 <http://example.org/se/0d59-9f0e-c05c-3bda>
2   rdf:type ex:ParkingSpot ;
(b) 3   dul:hasLocation ex:ACMEParkingPlace1 ;
4     ssn:attachedSystem <http://sensorportal.tu-braunschweig.de/ontology/spitfire.owl#sensor1> .
5
6 <http://sensorportal.tu-braunschweig.de/ontology/spitfire.owl#sensor1>
7   rdf:type ssn:Sensor ;
8   ssn:observes se:Occupancy ;
9   dul:isPartOf <http://example.org/se/0d59-9f0e-c05c-3bda> .

```

Figure 4. a) Example SPARQL query to find sensors at a given parking spot; b) example set of RDF triples describing the semantic entity and its virtual occupancy sensor. The *ParkingSpot* class is a subclass of *se:SemanticEntity*.

tic entity for a parking spot, one would query for sensors (e.g., magnetometers to detect the presence of metal) located at this parking spot. Additionally, the output of these sensors has to be mapped to the high-level states “occupied” or “free.” For that, we assume the existence of a RESTful web service that computes the high-level state from raw sensor output. We call such a web service a *virtual sensor*. Using a PUT operation, the virtual sensor can be configured to accept input from the previously found real sensors, and the computed high-level state can be retrieved by a GET operation from the virtual sensor in the same way as retrieving the output of a real sensor.

The newly created semantic entity is represented by a set of RDF triples describing its type and location, as well as providing pointers to the virtual sensor(s) computing its high-level state(s). Those virtual sensors in turn are also described by RDF triples so that they can be found by SPARQL queries.

Initially, virtual sensors have to be implemented manually, but over time a large set of reusable virtual sensors for common high-level states will be created and can be reused by posing SPARQL queries to find them. The virtual sensors either live on physical sensor nodes (avoiding the transmission of raw sensor time series across the Internet) or on Internet servers (if state computation exceeds the resources of a sensor node). If the real-world entity being represented by the semantic entity is covered by multiple sensors in a redundant fashion, the virtual sensor may also disable redundant sensors to save energy.

Figure 4a shows a SPARQL query to find all sensors located at parking spot “ACMEParkingPlace1.” Figure 4b shows the resulting set of RDF triples describing the semantic entity created for this parking spot. The reference to the newly created virtual occupancy sensor, the RDF description of which is shown further below, contains a link back to the semantic entity.

SEARCHING THINGS BY THEIR CURRENT STATES

An important functionality required by the use case sketched earlier is searching for semantic entities that exhibit a certain high-level state at the time of the query. As the high-level states

typically change very frequently, SPARQL is not directly applicable. Even CoAP’s observe function to report changes of a resource rather than streaming sensor values would not help as all sensors in the world would then push updates at high frequencies, which does not scale to an Internet of Things. Even installing an observation filter such that only values sought by a query are reported does not help, as then all relevant sensors would have to be tasked for each short query. We address this challenge by developing heuristics to efficiently identify entities that are *likely* to match a given search. We retrieve actual current state only for those likely candidates to check if they match [10].

We employ prediction models to compute the probability that the current high-level state of a semantic entity matches the value for which we are searching (e.g., parking spots that are empty). These prediction models are created by the virtual sensors (this is especially efficient if the virtual sensor lives on the physical sensor node that generates the input sensor data for the virtual sensor) and are periodically indexed by a search engine. When a search is performed, the search engine executes the indexed prediction models to obtain the matching probability without communicating with the virtual sensor. The semantic entities are then sorted by decreasing probability and contacted in this order to fetch their actual current state until enough matching entities have been found. We integrate this approach into SPARQL by encoding the prediction models as RDF triples and evaluating them in the SPARQL query. Using the ORDER BY clause, entities are then sorted by decreasing probability.

We employ different types of prediction models. The first type of model exploits periodic patterns in past states (e.g., a meeting room is occupied every Monday from 8 to 10). A second type of model exploits correlations of sensors (e.g., parking spots close to the entrance of a building are often all occupied, whereas spots further away are often free), so by knowing the occupancy of a parking spot one can infer the state of a neighboring parking spot with high probability. The prediction models we consider are based on simple statistics, and computational overhead for creating and executing them is small.

We employ prediction models to compute the probability that the current high-level state of a semantic entity matches the value we are searching for. These prediction models are created by the virtual sensors and are periodically indexed by a search engine.

```

1 <http://sensorportal.tu-braunschweig.de/ontology/spitfire.owl#sensor1>
2   rdf:type ssn:Sensor ;
(a) 3   ssn:observes se:Occupancy
4     dul:isPartOf <http://example.org/se/0d59-9f0e-c05c-3bda> ;
5     ex:FreeStateHasProbability "0.8" ;
6     ex:OccupiedStateHasProbability "0.2" .

1 SELECT DISTINCT ?entity
2 WHERE {
(b) 3   ?entity a ex:ParkingSpot ;
4       ssn:attachedSystem ?node .
5   ?node a ssn:Sensor ;
6       ex:FreeStateHasProbability ?prob .
7 } ORDER BY ?prob

```

Figure 5. a) RDF encoding of a simple prediction model for a virtual sensor; b) SPARQL query to compute a ranked list of parking spot semantic entities.

Figure 5a shows an extended version of the occupancy virtual sensor entity of Fig. 4b, encoding a very simple prediction model (i.e., the spot was occupied with probability 0.2 and free with probability 0.8 in the past). The SPARQL query in Fig. 5b returns a list of parking-spot semantic entities sorted by decreasing probability of being free according to the above prediction model.

REFERENCE IMPLEMENTATION

To demonstrate the feasibility of the proposed architecture, we realized a prototype using the Jena Semantic Web Framework. Figure 6 shows an overview of the architecture along with the steps involved in performing a search operation.

An RDF triple store (i.e., a distributed or grid/cloud-based database storing RDF triples) forms the core of the system. A crawler periodically scans the Semantic Web of Things for semantic entities and sensors, downloads metadata and prediction models using their RESTful web interfaces, converts this information into RDF triples (if necessary), and stores them in the triple store. iSense sensor nodes equipped with motion and temperature sensors are deployed in our office building experiment. Their raw output is converted into a high-level state reflecting room occupancy by a virtual sensor. The sensor nodes run 6LoWPAN/IPv6 on the network layer and an implementation of the current IETF CoAP draft on the transport layer to provide a RESTful interface to the outside world. A proxy translates between CoAP/UDP-based service invocations and RESTful HTTP/TCP web service calls that are being issued by the crawler or arbitrary client applications. In addition to physical sensors, our system supports so-called web sensors, that is, sensors whose output is published on web pages that are parsed by the crawler to extract sensor values. In particular, we are using a city-wide online parking web application to obtain real-time information about parking lot occupancy throughout the city.

Semantic data stored in the triple store can be queried using a SPARQL query engine. Our aim is to also support non-expert users who would be unable to formulate SPARQL queries. We therefore offer a simple graphical user front-end that accepts a keyword-based query and converts it into a SPARQL query at the cost of reduced query expressiveness. A map-based interface shows locations of the matching semantic entities.

The system prototype is operational and implements the use case described earlier. The system accesses about 40 physical sensors (temperature, movement detectors) attached to 20 sensor nodes in our office building, as well as 25 web sensors (parking lot occupancy). We have also studied the efficiency of the underlying content-based search algorithm in [10], showing that the communication overhead is reduced by up to one order of magnitude compared to a pull-based baseline.

Future steps include (distributed/grid/cloud-based) triple store scalability for sensor data, which uses domain knowledge about the stored triples and specialized data placement policies to optimally distribute triples.

Also, we are working on supporting better reasoning by making use of LOD ontologies in queries. Finally, we are currently working on integrating a service for semi-automatic sensor annotation into the system. The reference implementation will be provided at <http://spitfire-project.eu>.

CONCLUSIONS

Existing semantic sensor web technologies enable the integration of sensors into the web, but the underlying model is focused on sensors rather than on things and their high-level states. Existing approaches also lack integration with the LOD cloud, a quickly growing and open base of semantic world knowledge. With SPITFIRE, we work toward a Semantic Web of Things by providing abstractions for things, fundamental services for search and annotation, as well as integrating sensors and things into the LOD cloud. Using linked data principles makes sensor data easily accessible for applications via existing mechanisms deployed on the web, which will significantly speed up the uptake of IoT technologies. We demonstrate the feasibility by an operational prototype used to realize a representative use case. As SPITFIRE is in line with the most promising developments on both the IoT and web side, the system will benefit from any advances made in these areas implicitly.

As it was difficult to foresee the wealth of current web applications back when the web was created, we have to wait and see how people will use the Semantic Web of Things. It is also hard to predict if a Semantic Web of Things will be as broadly adopted as the web is today. One indicator is that LOD has already achieved significant uptake by governments (including those of the United Kingdom and United States), the media sector (BBC), life sciences, geo information systems, and web companies (Freebase). Making sensor data part of this data pool is clearly beneficial, as then integration with knowledge from arbitrary sources is possible. For example, sensors and their data can be linked to geographic data (correlated natural phenomena), user-generated data (social feedback), government data (census information), life-science data (causes and effects of diseases), and so on.

A strong indicator of whether this line of development will be successful in the long run is also provided by the exponentially growing amount of linked data and the support by major

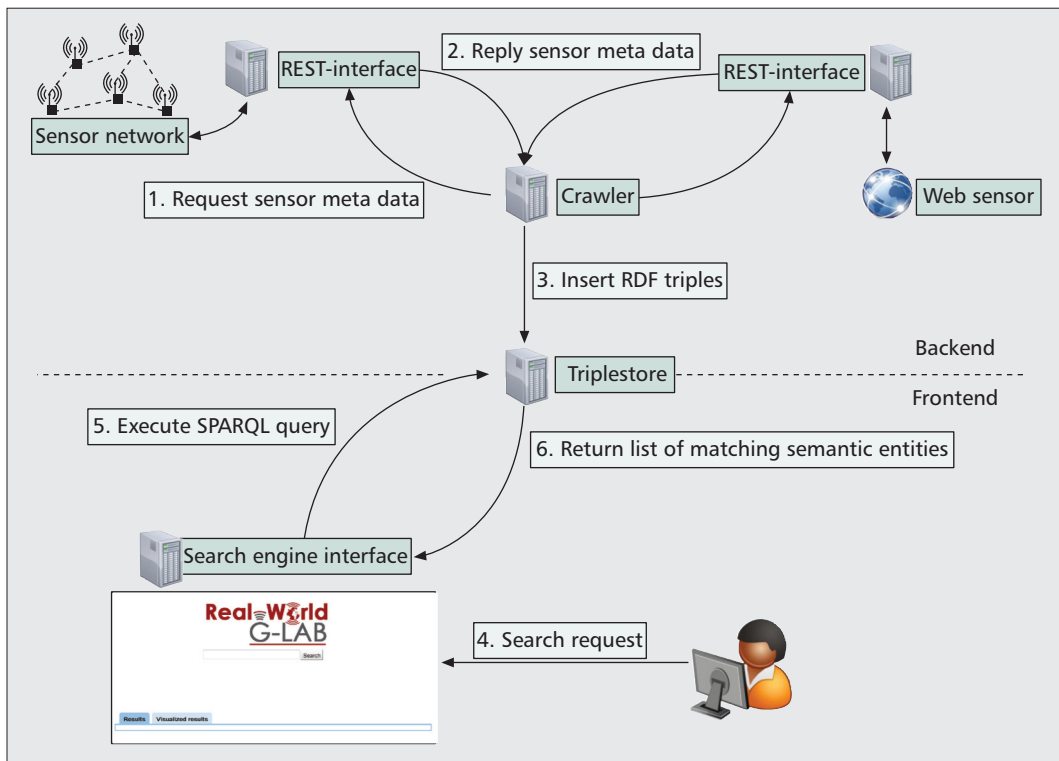


Figure 6. Architecture of the reference implementation.

players. Since its beginnings in 2007, the LOD cloud has grown from 12 datasets to 203 datasets in 2010 with over 25 billion triples interlinked with 395 million links. Industry initiatives such as Google's Rich Snippet (2009), Facebook's Open Graph (2010), and very recently Schema.org (2011) all aim at adding semantic markups to web pages to improve search and discovery capabilities for the end user, and also confirm its uptake at web scale.

REFERENCES

- [1] Open Geospatial Consortium, "Sensor Web Enablement (SWE)," 2010.
- [2] A. Sheth, C. Henson, and S. Sahoo, "Semantic Sensor Web," *IEEE Internet Computing*, July/Aug. 2008, 2008, pp. 78–83.
- [3] G. Montenegro et al., "Transmission of IPv6 Packets over IEEE 802.15.4 Networks," IETF RFC 4944, Sept. 2007; <http://www.ietf.org/rfc/rfc4944.txt>
- [4] Z. Shelby et al., "Constrained Application Protocol (CoAP) (CoRE Working Group)," <http://www.ietf.org/id/draft-ietf-core-coap-06.txt> (07/07/2011), 2011.
- [5] R. Fielding, *Architectural Styles and the Design of Network-Based Software Architectures*, Ph.D. dissertation, UC Irvine, 2000.
- [6] Z. Shelby, "Embedded Web Services," *IEEE Wireless Commun.*, vol. 17, no. 6, Dec. 2010, pp. 52–57; <http://dx.doi.org/10.1109/MWC.2010.5675778>.
- [7] T. Lee et al., "The Semantic Web," *Sci. Amer.*, vol. 284, no. 5, 2001, pp. 34–43.
- [8] S. Jirka, A. Bröring, and C. Stasch, "Discovery Mechanisms for Sensor Web," *Sensors*, 2009.
- [9] J. Pschorr et al., "Sensor Discovery on Linked Data," Kno.e.sis, tech. rep., 2010.
- [10] B. M. Elahi et al., "Sensor Ranking: A Primitive for Efficient Content-Based Sensor Search," *Proc. 2009 Intl. Conf. Info. Processing in Sensor Networks*, 2009, pp. 217–28.

BIOGRAPHIES

DENNIS PFISTERER is a senior researcher and tenured lecturer at the University of Lübeck, Germany. He was a research assistant at the European Media Laboratory (<http://www.>

eml.org) in Heidelberg in the area of resource adaptive systems and at the Braunschweig Institute of Technology where he worked on sensor networks. In 2005 he went to Lübeck to continue his research at the Institute of Telematics. After his Ph.D., he broadened his research interests and now works on sensor networks in general, their integration with the future Internet, as well as service oriented architecture (SOA).

KAY RÖMER is an associate professor at the Institute for Computer Engineering of the University of Lübeck, Germany. He holds a Diploma degree (1999) in informatics from the University of Frankfurt/Main, Germany, and a doctoral degree (2005) in informatics from ETH Zurich, Switzerland. Between 2005 and 2009 he was a senior researcher at ETH Zurich. His research interests encompass networked embedded systems, in particular wireless sensor networks and the Internet of Things.

DANIEL BIMSCHAS is an assistant researcher at the Institute of Telematics at the University of Lübeck. He received his Master's degree in computer science from the University of Lübeck in 2009. His research is focused on experimental research platforms for the future Internet, the Internet of Things (IoT), as well as on integration of and application development for IoT devices.

HENNING HASEMANN received his Diploma in computer science at Braunschweig University of Technology, Germany (2010). He is currently a Ph.D. student in the Algorithms Group of Prof. Fekete at Braunschweig University of Technology, Germany. His main research interests are theoretical and practical aspects of distributed algorithms for sensor networks and motion planning.

MANFRED HAUSWIRTH is vice-director of the Digital Enterprise Research Institute (DERI), Galway, Ireland, and a professor at the National University of Ireland, Galway (NUIG). He holds an M.S. (1994) and a Ph.D. (1999) in computer science from the Technical University of Vienna. He has published over 100 papers on semantic sensor networks and middleware, large-scale semantics-enabled distributed information systems, and other domains. He has co-authored a book on distributed software architectures and several book chapters.

MARCEL KARNSTEDT finished his Ph.D. in 2009 at the Databases and Information Systems Group at Ilmenau University of

As it was difficult to foresee the wealth of current web applications back when the web was created, we have to wait and see how people will use the Semantic Web of Things. It is also hard to predict if a Semantic Web of Things will be as broadly adopted as the web is today.

Technology, Germany. Since 2009 he has been working as a postdoctoral researcher at DERI, NUIG. His research interests are in large-scale data management, query processing, and data mining. He has more than 50 publications in areas like P2P databases, stream and graph mining, and distributed query processing. More information can be found at <http://www.marcel.karnstedt.com>.

OLIVER KLEINE is an assistant researcher at the Institute of Telematics at the University of Lübeck. He received his Master's degree in computer science from the University of Lübeck in 2010. His research interests are focused on concepts and application development for the Internet of Things.

ALEXANDER KRÖLLER is a senior researcher and lecturer at TU Braunschweig. He received his Diploma in mathematics from TU Berlin and his Ph.D. from TU Braunschweig; his thesis, *Algorithmic Methods for Wireless Sensor Networks*, was nominated for the Dissertation Award of the German Society of Computer Science. He specializes in algorithmic methods for distributed systems such as wireless sensor networks. He is a member of the Steering Committee of Wireless Sensor Networks: Theory and Practice.

MYRIAM LEGGIERI is a Ph.D. student at DERI, interested in linked sensor data modeling. She has previously worked on linked sensor data applied to collaborative software development for a project under the Google Summer of Code 2009 program, with the University of Hawaii. She has a Master's degree on the same topic from the University of Bari, Italy.

Richard Mietz received his diploma degree in computer science in 2009 from the University of Lübeck. Since 2009 on he is a Ph.D. student in the Working Group of Prof. Römer at the Institute of Computer Engineering, University of Lübeck. He is working in the BMBF-funded Real-World G-

Lab project on the intersection of wireless sensor networks and the Internet of Things.

MAX PAGEL is a Ph.D. student at Braunschweig University of Technology. After acquiring his Master's degree in the highly interdisciplinary course of studies Computational Life Science at the University of Lübeck, he is now working in the field of sensor networks and the Internet of Things. His special interests lie in algorithms for sensor networks, the efficient integration of sensors with the future Internet, and network data analysis.

ALEXANDRE PASSANT received his M.Sc. and M.Eng from Université-Paris Dauphine and his Ph.D. in 2009 from Université-Paris Sorbonne. From 2008 to 2011 he worked as a postdoctoral researcher at DERI. From 2011 on he is a research fellow and leader of the Social Software Unit. From 2009 on he has also been an adjunct lecturer at the National University of Ireland. His research interests are semantic web and knowledge representation.

RAY RICHARDSON is head of the research and development team at Daysha Consulting. He obtained his Ph.D. in computer science from Dublin City University in 1995. As a researcher at Bell Laboratories he led a team investigating the use of semantic web service technology to maximize the efficiency of large supply chain networks. His research activities have focused on the structured collaboration activities. He leads the team responsible for building the Open Concurrent Design Server.

CUONG TRUONG received his B.Sc. IT degree from the National University of Vietnam, and his M.Sc. IT degree from Anna University, Madras, India. He is currently a Ph.D. candidate at the Institute of Computer Engineering, University of Lübeck. His research interests include coverage and routing in sensor networks, quality of service in sensor networks, and sensor data mining and clustering in the Internet of Things.